

Mini Project : House Price Prediction Model

Problem Statement : Build a Machine Learning model that predicts the price of a house based on its features using regression algorithms.

Tasks Performed :

- Load and clean the "**House Price India.csv**" dataset.
- Perform **basic analysis and visualization** (e.g., condition of house vs. average price).
- Train and compare three regression models:
 - Decision Tree Regressor (with hyperparameter tuning using GridSearchCV)
 - Linear Regression
 - Random Forest Regressor (with hyperparameter tuning using GridSearchCV)
- Evaluate models using:
 - **Mean Squared Error (MSE)**
 - **Mean Absolute Error (MAE)**
- Save the **best model** as model.pkl using joblib for future use (e.g., in a web app like Streamlit or Flask).

Results and Interpretation:

After training and evaluating the three models on the test set, we compare their performance based on **MSE** and **MAE**:

- **Decision Tree Regressor**
 - MSE: 67739876625.71646
 - MAE: 161090.87323793155
- **Linear Regression**
 - MSE: 65514176740.267784
 - MAE: 165026.2962930765
- **Random Forest Regressor**
 - MSE: 61319014579.76153
 - MAE: 158602.34822515058

Code :

App.py:

```
import streamlit as st
import numpy as np
import joblib

model = joblib.load('model1.pkl')

st.title("House Price Prediction")

st.divider()

st.write("Enter the details of the house to predict its price:")

st.divider()

bedrooms = st.number_input("Number of Bedrooms", min_value=1, value=1)
bathrooms = st.number_input("Number of Bathrooms", min_value=0, value=1)
livingarea = st.number_input("Square Footage of Living Area", min_value=0,
value=2000)
condition = st.number_input("Condition of the House", min_value=0, value=3)
noofschools = st.number_input("Number of Schools Nearby", min_value=0,
value=0)

st.divider()

X = [[bedrooms, bathrooms, livingarea, condition, noofschools]]

predictbutton = st.button("Predict Price")

if predictbutton:
    st.balloons()

    X_array = np.array(X)
    prediction = model.predict(X_array)

    st.write(f"⚠️ The predicted price of the house
is: {prediction[0]:,.2f}")

else:
    st.write("Please click the **Predict Price** button after entering
details.")

# Command to run : streamlit run g:/housePricePrediction/app1.py
```

Main.py:

```
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor

from sklearn.metrics import mean_squared_error, mean_absolute_error

import joblib

data = pd.read_csv('House Price India.csv')

# Show basic info about columns, data types, and non-null counts
print("\n--- Data Info ---")
print(data.info())

#Basic Statistics of Target Variable (Price)
# Describe the "Price" column: count, mean, std, min, quartiles, max
stats = data["Price"].describe().reset_index()

# Round the statistics values to 2 decimal places for better readability
stats["Price"] = stats["Price"].round(2)

print("\n--- Price Statistics ---")
print(stats)

#Check Missing and Duplicate Values

# Total number of missing values in the entire DataFrame
total_missing = data.isna().sum().sum()
print(f"\nTotal missing values in dataset: {total_missing}")

# Total number of fully duplicated rows
total_duplicates = data.duplicated().sum()
print(f"Total duplicated rows in dataset: {total_duplicates}")

#Handle Missing and Duplicate Values

# Drop rows with any missing value
data.dropna(inplace=True)

# Drop exactly duplicated rows
data.drop_duplicates(inplace=True)

print("\n--- Data after cleaning (head) ---")
```

```

print(data.head())

data.groupby("condition of the
house")["Price"].mean().sort_values(ascending=True).plot(kind='bar')
plt.title("Condition of the House vs Average Price")
plt.ylabel("Average Price")
plt.xlabel("Condition of the House")
plt.show()

X = data[['number of bedrooms','number of bathrooms','living area','condition
of the house','Number of schools nearby']]

X

y = data["Price"]

y

X_train, X_test, y_train, y_test = train_test_split(X,
y,test_size=0.2,random_state=42)

print("\nTraining set size:", X_train.shape, y_train.shape)
print("Test set size:", X_test.shape, y_test.shape)

param_grid_tree = {
    "criterion": ["squared_error", "friedman_mse", "absolute_error"],
    "splitter": ["best", "random"],
    "max_depth": [None, 10, 20, 30, 40, 50],
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 4]
}

# Base Decision Tree model
tree_model = DecisionTreeRegressor(random_state=42)

# GridSearchCV tries all combinations of parameters and picks the best
grid_tree = GridSearchCV(
    estimator=tree_model,
    param_grid=param_grid_tree,
    scoring="neg_mean_squared_error", # we want to minimize
MSE                                # use all CPU cores
)

# Fit the model on training data
grid_tree.fit(X_train, y_train)

print("\n--- Best Parameters for Decision Tree ---")

```

```

print(grid_tree.best_params_)

# Predictions on test set using best Decision Tree model
tree_preds = grid_tree.predict(X_test)

# Evaluate Decision Tree
tree_mse = mean_squared_error(y_test, tree_preds)
tree_mae = mean_absolute_error(y_test, tree_preds)

print("\nDecision Tree Performance:")
print("MSE:", tree_mse)
print("MAE:", tree_mae)

# Initialize Linear Regression model
lr = LinearRegression()

# Train on training data
lr.fit(X_train, y_train)

# Predict on test data
lr_preds = lr.predict(X_test)

# Evaluate Linear Regression
lr_mse = mean_squared_error(y_test, lr_preds)
lr_mae = mean_absolute_error(y_test, lr_preds)

print("\nLinear Regression Performance:")
print("MSE:", lr_mse)
print("MAE:", lr_mae)

# Base Random Forest model
rf_model = RandomForestRegressor(random_state=42)

# Parameter grid for Random Forest
param_grid_rf = {
    "max_depth": [5, 10, 15],
    "n_estimators": [2, 3, 4, 5, 6, 7, 8, 9, 10]
}

# GridSearchCV for Random Forest
grid_rf = GridSearchCV(
    estimator=rf_model,
    param_grid=param_grid_rf,
    scoring="neg_mean_squared_error"
)

# Fit Random Forest with GridSearch on training data

```

```

grid_rf.fit(X_train, y_train)

print("\n--- Best Parameters for Random Forest ---")
print(grid_rf.best_params_)

# Predictions with best Random Forest model
rf_preds = grid_rf.predict(X_test)

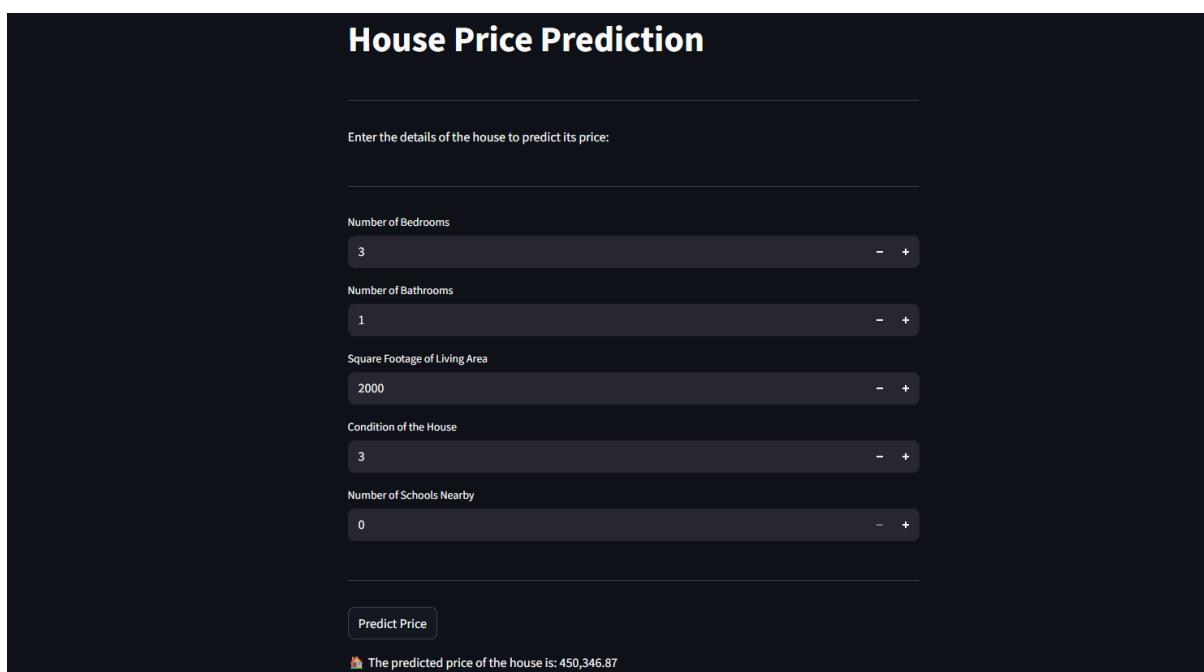
# Evaluate Random Forest
rf_mse = mean_squared_error(y_test, rf_preds)
rf_mae = mean_absolute_error(y_test, rf_preds)

print("\nRandom Forest Performance:")
print("MSE:", rf_mse)
print("MAE:", rf_mae)

#Save best model as pkl file
joblib.dump(grid_rf, 'model1.pkl')

```

Output :



The screenshot shows a web-based application titled "House Price Prediction". The interface is dark-themed with white text and light-colored input fields. At the top, it says "Enter the details of the house to predict its price:". Below this, there are five input fields for "Number of Bedrooms" (set to 3), "Number of Bathrooms" (set to 1), "Square Footage of Living Area" (set to 2000), "Condition of the House" (set to 3), and "Number of Schools Nearby" (set to 0). Each input field has a minus and plus button on the right. At the bottom, there is a "Predict Price" button and a message: "The predicted price of the house is: 450,346.87".

House Price Prediction

Enter the details of the house to predict its price:

Number of Bedrooms

4

-

+

Number of Bathrooms

1

-

+

Square Footage of Living Area

2000

-

+

Condition of the House

3

-

+

Number of Schools Nearby

0

-

+

Predict Price

 The predicted price of the house is: 444,468.62

House Price Prediction

Enter the details of the house to predict its price:

Number of Bedrooms

3

-

+

Number of Bathrooms

1

-

+

Square Footage of Living Area

2000

-

+

Condition of the House

5

-

+

Number of Schools Nearby

1

-

+

Predict Price

 The predicted price of the house is: 500,989.10